

AI Research Agent — Research Quality & Integrity Guide

Research integrity checklist

The application should never claim that an answer is guaranteed to be correct. Its defensible promise is evidence-grounded research with transparent sources, source-quality prioritization, explicit uncertainty, and references.

Source handling

1. Discover sources
2. Deduplicate URLs
3. Rank source domains
4. Retrieve original pages
5. Extract readable evidence
6. Provide evidence to the research agent
7. Generate report with numbered references

Model behavior rules

Never invent sources, URLs, statistics, dates, quotations or findings. Prefer primary and authoritative sources. Distinguish facts from interpretation. Preserve dates and context. Report disagreements. Do not use model memory as a replacement for retrieved evidence.

High-stakes topics

Legal, medical, financial, regulatory and similarly high-stakes research requires additional human verification. The application should not present itself as a substitute for a qualified professional.

Future upgrades

Recommended next upgrades are a citation auditor, claim-to-evidence mapping, publication-date extraction, domain-specific source allowlists, retry/backoff, rate-limit handling, structured report export, source caching, research history, and automated evaluation tests.